



RetiGate Sensing Overhead by Pipeline Stage



Sensing vs Inference Latency (separate hardware, pipelined)



Sensing: Apple M3 Pro CPU (n=500 frames)
 Inference: NVIDIA Tesla T4 GPU (CUDA events, n=200)
 Hardware is separate — latencies are NOT additive